



# Scalable Approximation of Quantitative Information Flow in Programs


Fabrizio Biondi, Michael A. Enescu, Annelie Heuser, Axel Legay, Kuldeep S.Meel, Jean Quilbeuf

EQUIPE TAMIS

CENTRE DE RENNES -  
BRETAGNE ATLANTIQUE

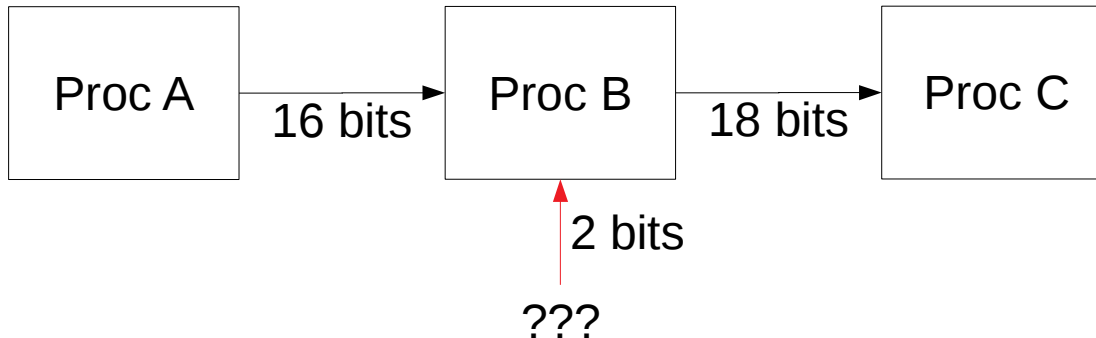
7 Jan 2018

# Information Flow Quantification



Information flow from A to B =  $\log_2(\text{possible values}(x))$  bits

8 values  $\rightarrow$  3 bits      64 values  $\rightarrow$  6 bits  
4294967296 values  $\rightarrow$  32 bits



# Heartbleed leaks memory...

```
int dtls1_process_heartbeat(SSL *s) {  
  
    unsigned char *p = &s->s3->rrec.data[0], *pl;  
    unsigned short hbtype;  
    unsigned int payload;  
    unsigned int padding = 16;  
    //...  
    hbtype = *p++;  
    n2s(p, payload);
```

```
    if (1+2 + payload+16 > s->s3->rrec.length)  
        return 0; /* missing in bugged version */
```

← This is the fix

```
    if (hbtype == TLS1_HB_REQUEST) {  
        unsigned char *buffer, *bp;  
        unsigned int write_length =  
            1 + 2 + payload + padding;  
        //..
```

```
        buffer = OPENSSL_malloc(write_length);  
        bp = buffer;  
        *bp++ = TLS1_HB_RESPONSE;  
        s2n(payload, bp);  
        memcpy(bp, pl, payload);  
        //send buffer ...
```

← Will send kernel memory

```
    }  
}
```

# ...can we detect it?

```
int dtls1_process_heartbeat(SSL *s) {  
  
    unsigned char *p = &s->s3->rrec.data[0], *pl;  
    unsigned short hbtype;  
    unsigned int payload;  
    unsigned int padding = 16;  
    //...  
    hbtype = *p++;  
    n2s(p, payload);  
  
    if (1+2 + payload+16 > s->s3->rrec.length)  
        return 0; /* missing in bugged version */  
  
    if (hbtype == TLS1_HB_REQUEST) {  
        unsigned char *buffer, *bp;  
        unsigned int write_length =  
            1 + 2 + payload + padding;  
        //..  
        buffer = OPENSSL_malloc(write_length);  
        bp = buffer;  
        *bp++ = TLS1_HB_RESPONSE;  
        s2n(payload, bp);  
        memcpy(bp, pl, payload);  
        //send buffer ...  
    }  
}
```

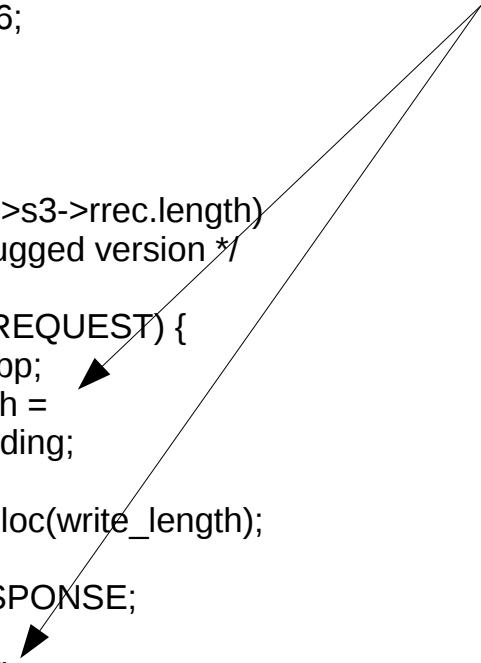
## Approach\*:

- 1) get SAT constraints on variables you care about
- 2) use projected SAT model counter to count possible values

\* extremely simplified, details in paper

# ...can we detect it?

```
int dtls1_process_heartbeat(SSL *s) {  
  
    unsigned char *p = &s->s3->rrec.data[0], *pl;  
    unsigned short hbtype;  
    unsigned int payload;  
    unsigned int padding = 16;  
    //...  
    hbtype = *p++;  
    n2s(p, payload);  
  
    if (1+2 + payload+16 > s->s3->rrec.length)  
        return 0; /* missing in bugged version */  
  
    if (hbtype == TLS1_HB_REQUEST) {  
        unsigned char *buffer, *bp;  
        unsigned int write_length =  
            1 + 2 + payload + padding;  
        //..  
        buffer = OPENSSL_malloc(write_length);  
        bp = buffer;  
        *bp++ = TLS1_HB_RESPONSE;  
        s2n(payload, bp);  
        memcpy(bp, pl, payload);  
        //send buffer ...  
    }  
}
```



## Approach\*:

- 1) get SAT constraints on variables you care about
- 2) use projected SAT model counter to count possible values

## Result:

- Generated SAT formula with 39272 clauses in <1s
- ... model counter timeout :(

\* extremely simplified, details in paper

# Solution: approximate SAT counting

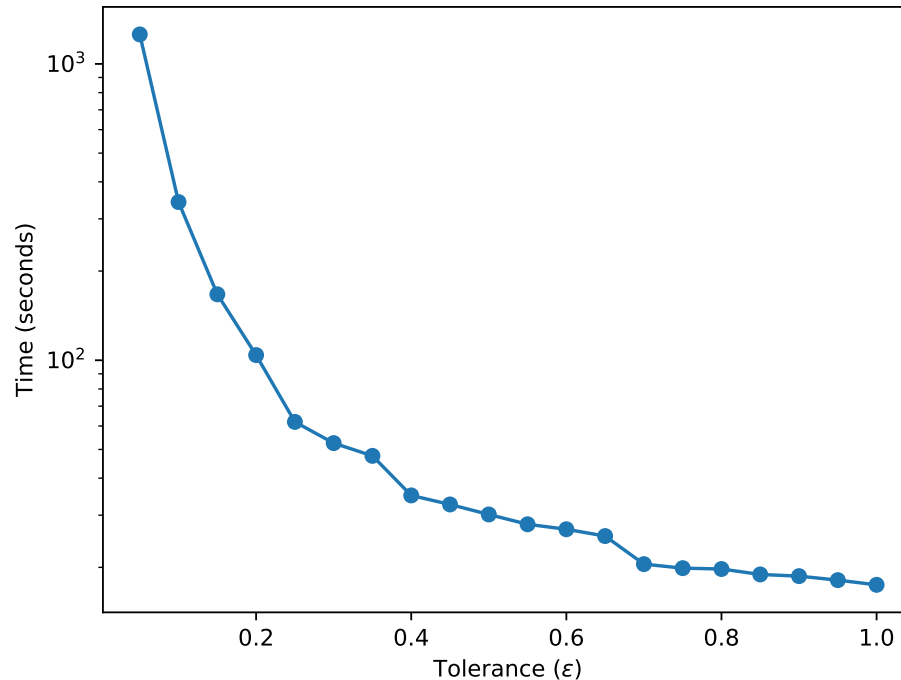
- Counting precision not very important (we are taking a log anyway)
- Can we trade precision for speed by approximating?
- Klebanov et al.<sup>1</sup> already did it!
- Except...
  - We couldn't reproduce their results (authors were nice and helpful)
  - Error in Theorem 2.12 overestimates termination probability
  - Error in Theorem 2.6 overestimates precision (details in paper)
- So we used ApproxMC instead (v2 has projected model counting<sup>2</sup>)

<sup>1</sup>V. Klebanov, A. Weigl, and J. Weisbarth. Sound probabilistic #SAT with projection. In QAPL 2016

<sup>2</sup>S. Chakraborty, K. S. Meel, and M. Y. Vardi. Algorithmic improvements in approximate counting for probabilistic inference: From linear to logarithmic SAT calls. In IJCAI 2016

# Performance-precision trade-off

Preprocessed AppleTalk benchmark:



Tolerance of 1: approximation is from 0.5 to 2 times the real value  
→ exactly +/- 1 bit of information

# Benchmarks

Benchmarks from [24, 5, 23, 19]

Experiment name	sharpCDCL leakage	ApproxMC2 leakage	Relative error	sharpCDCL time	ApproxMC2 time	Speedup factor
e-purse	5.00	5.00	0%	0.06	0.28	-4.67
pw-checker	1.00	1.00	0%	0.00	0.00	—
sum-query	>22.49	32.00	*	t/o	0.87	*
10random	3.32	3.32	0%	0.00	0.00	—
bsearch16	16.00	16.00	0%	3.40	0.49	6.90
bsearch32	>22.87	32.00	*	t/o	2.13	*
mix-dupl	16.00	16.00	0%	5.91	0.20	29.60
sum32	>22.48	32.00	*	t/o	0.89	*
illustr.	4.09	4.09	0%	0.00	0.01	—
mask-cpy	16.00	16.00	0%	6.02	0.20	30.1
sanity-1	>22.82	31.04	*	t/o	0.94	*
sanity-2	>22.92	31.00	*	t/o	1.07	*
check-cpy	>22.51	32.00	*	t/o	0.88	*
copy	>22.49	32.00	*	t/o	0.84	*
div-by-2	22.79	31.00	*	t/o	1.06	*
implicit	>2.81	2.81	0%	0.00	0.01	—
mul-by-2	>22.46	31.00	*	t/o	0.89	*
popcnt	5.04	5.04	0%	0.00	0.01	—
simp-mask	8.00	8.00	0%	0.00	0.05	—
switch	4.25	4.25	0%	0.00	0.00	—
tbl-lookup	>22.45	32.00	*	t/o	0.88	*



# Benchmarks

Benchmarks from [24, 5, 23, 19]

Experiment name	sharpCDCL leakage	ApproxMC2 leakage	Relative error	sharpCDCL time	ApproxMC2 time	Speedup factor
ddp	error	128.00	*	error	23.50	*
ddp.pp	error	128.00	*	error	19.55	*
popcount	<b>5.04</b>	<b>5.04</b>	<b>0%</b>	<b>0.00</b>	<b>0.01</b>	—
sanitize	<b>4.00</b>	<b>4.00</b>	<b>0%</b>	<b>0.00</b>	<b>0.00</b>	—
openssl.1	<b>8.00</b>	<b>8.00</b>	<b>0%</b>	<b>1.44</b>	<b>70.66</b>	<b>-49.10</b>
openssl.2	<b>16.00</b>	<b>16.00</b>	<b>0%</b>	<b>4.63</b>	<b>75.39</b>	<b>-16.30</b>
openssl.3	>22.24	24.00	*	t/o	92.47	*
openssl.4	>22.91	32.00	*	t/o	86.32	*
openssl.5	>23.10	40.00	*	t/o	87.74	*
openssl.6	error	48.00	*	error	89.60	*
openssl.7	error	56.00	*	error	91.98	*
openssl.8	error	64.00	*	error	98.04	*
openssl.9	error	72.00	*	error	97.41	*
openssl.10	error	80.00	*	error	112.71	*
openssl.15	error	t/o	*	error	t/o	*
openssl.20	error	160.00	*	error	142.48	*
swirl	>12.82	t/o	*	t/o	t/o	—
10random	<b>3.32</b>	<b>3.32</b>	<b>0%</b>	<b>0.00</b>	<b>0.01</b>	—
bsearch16	<b>16.00</b>	<b>16.00</b>	<b>0%</b>	<b>4.16</b>	<b>0.68</b>	<b>6.12</b>
bsearch16.pp	<b>16.00</b>	<b>16.00</b>	<b>0%</b>	<b>3.73</b>	<b>0.35</b>	<b>10.70</b>
bsearch32	>22.79	32.00	*	t/o	3.21	*
bsearch32.pp	>22.90	32.00	*	t/o	6.93	*
fx	<b>16.00</b>	<b>16.00</b>	<b>0%</b>	<b>5753.42</b>	<b>7307.61</b>	<b>-1.27</b>
mixdup	<b>16.00</b>	<b>16.00</b>	<b>0%</b>	<b>8.44</b>	<b>0.22</b>	<b>38.40</b>
sum.32	>22.78	32.00	*	t/o	0.98	*

# So how about Heartbleed?

```
int dtls1_process_heartbeat(SSL *s) {  
  
    unsigned char *p = &s->s3->rrec.data[0], *pl;  
    unsigned short hbtype;  
    unsigned int payload;  
    unsigned int padding = 16;  
    //...  
    hbtype = *p++;  
    n2s(p, payload);  
  
    if (1+2 + payload+16 > s->s3->rrec.length)  
        return 0; /* missing in bugged version */  
  
    if (hbtype == TLS1_HB_REQUEST) {  
        unsigned char *buffer, *bp;  
        unsigned int write_length =  
            1 + 2 + payload + padding;  
        //..  
        buffer = OPENSSL_malloc(write_length);  
        bp = buffer;  
        *bp++ = TLS1_HB_RESPONSE;  
        s2n(payload, bp);  
        memcpy(bp, pl, payload);  
        //send buffer ...  
    }  
}
```

## Approach\*:

- 1) get SAT constraints on variables you care about
- 2) use *approximate* projected SAT model counter to *estimate* possible values

## Result:

- Generated SAT formula with 39272 clauses in <1s
- Computed flow of ~15 bytes in 25s
- Reducing confidence gives ~15.1 bytes in 2s
- Normal flow should be 1 byte\*
- **Bug found!**

\* extremely simplified, details in paper

# Conclusions

- Information flow quantification can detect interesting bugs
- Approximate quantification is significantly faster than precise
  - Large performance increase for negligible precision loss
- Approximate quantification scales to real-world code and bugs
  - Modeling still complex (but mostly engineering problem)
- Future work:
  - Lower-bound estimation is sufficient and faster
  - Multiple upwards/downwards passes for refinement in large programs
  - ...mostly engineering?

Thank you for your attention!